

ECE 220

Lectures 03: Repeated Code
TRAPs and Subroutines

Slides by Sayan Mitra,, Sanjay Patel, Yih-Chun Hu

Outline

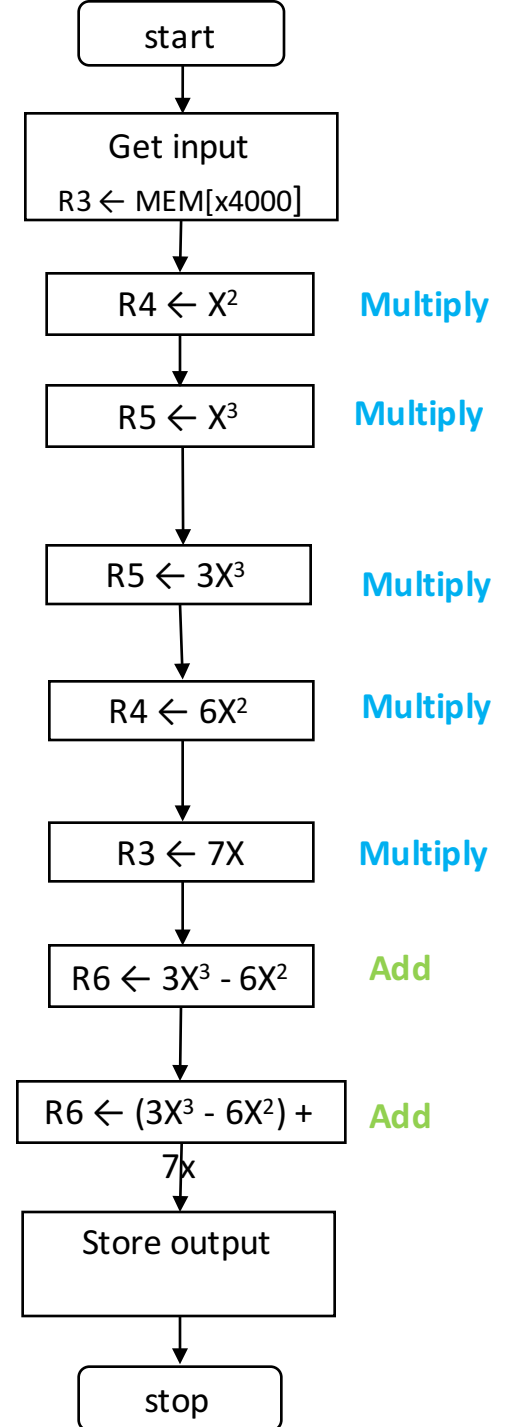
- Chapter 9
- Repeated code: TRAPs and Subroutines
- Key concepts
 - Lookup table: for starting address of subroutines/TRAPS (vector table)
 - Preserving register and PC values
- Instructions
 - TRAP
 - RET
 - JSR, JSRR

Observation 1

Example problem: Compute $y = 3x^3 - 6x^2 + 7x$ for any input $x > 0$

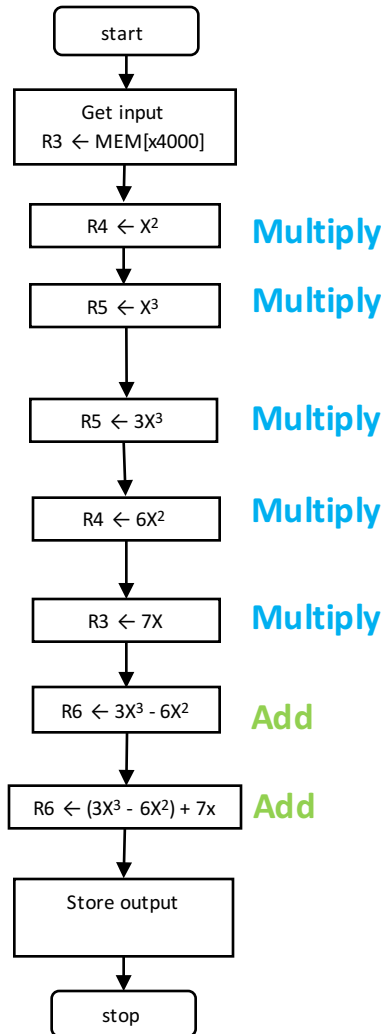
Programs have lots of **repetitive** code fragments

```
; multiply R0 ← R1 * R2
MULT   AND R0, R0, #0      ; R0 = 0
LOOP   ADD R0, R0, R2      ; R0 = R0 + R2
        ADD R1, R1, #-1    ; decrease counter
        BRp LOOP
```



Implementation Option 1

Issues ?



```
;; LC-3 Assembly Program
.ORIG x3000
LDI R3, Xaddr;  R3 ← x
ADD R1, R3, #0;
; Multiply R4 ← R1 * R3 x2
...
...
; Multiply R5 ← R4 * R3 x3
...
; Multiply R5 ← R5 * 3 (3x3)
...
; Multiply R4 ← 6 * R4
```

*Sometimes programs are compiled to look like this for better performance: inline functions

3 virtues of great programmer ---

Larry Wall (Author of Perl)

- **Laziness:** write labor-saving programs that other people will find useful and document what you wrote so you don't have to answer so many questions about it.
- **Impatience:** write programs that don't just react to your needs, but actually anticipate them.
- **Hubris:** write programs that other people won't want to say bad things about.

Another example: Code (from last lecture) for reading characters from keyboard

```
START      LDI      R1,
KBSR_ADDR

          BRzp     START

          LDI      R0,
KBDR_ADDR


          BRnzp    NEXT_TASK
          ...
KBSR_ADDR  .FILL    xFE00
KBDR_ADDR  .FILL    xFE02
```

- Very common... would be replicated often
- Too many specific details for most programmers
 - know address of KBDR (xFE02) and KBSR (xFE00)
 - use the registers correctly (polling, data format)
- Improper usage could breach security of the system!
- Idea: Make this part of the (operating) system
- In general, provide “pieces of code” or *subroutines or system calls or service routine* that do all of these low-level tasks correctly
 - User **invokes or calls** subroutine
 - Operating system code performs operation
 - **Returns** control to user program

How to make this idea work?

User program **invokes or calls** subroutine; OS code performs operation; **Returns** control to user program

- The actual code of the service routine
- Mechanism for invocation
 - TRAP Instruction, e.g., TRAP x20
 - TRAP vector (8 bits)
 - How to find address service routine?



Address	Contents	Comments
x0000		;system space;
...		
x00FF		
x0400		; code for GETC
....		
x0430		; code for OUT
...		
x3000		; user program
...		
	TRAP x20	; call to
xFE00		; Device registers

TRAP Vector Table for LC3

vector	address	symbol	routine
...			
x20	x0400	GETC	read a single character (no echo)
x21	x0430	OUT	output a character to the monitor
x22	x0450	PUTS	write a string to the console
x23	x04A0	IN	print prompt to console, read and echo character from keyboard
X23	x04E0	PUTSP	write a string to the console; two chars per memory location
x25	xFD70	HALT	halt the program
...			

Look-up table decouples names of subroutines (GETC) from the location of its implementation in memory

1. exit	801d4a74	T
2. fork	801d7980	T
3. read	801eb584	T
4. write	801eb958	T
5. open	800b13a4	T
6. close	801ccab4	T
7. wait4	801d56bc	T
9. link	800b18e8	T
10. unlink	800b1ff0	T
12. chdir	800b0c60	T
13. fchdir	800b0af0	T
14. mknod	800b14bc	T
15. chmod	800b2b40	T
16. chown	800b2c9c	T
18. getfsstat	800b088c	T
20. getpid	801dc20c	T
23. setuid	801dc4c0	T
24. getuid	801dc290	T
25. geteuid	801dc2a0	T
26. ptrace	801e812c	T
27. recvmsg	8020a8fc	T
28. sendmsg	8020a444	T
29. recvfrom	8020a528	T
30. accept	80209dfc	T
31. getpeername	8020abc8	T
32. getsockname	8020ab18	T
33. access	800b24ac	T
34. chflags	800b2928	T
35. fchflags	800b29f0	T
36. sync	800b0320	T
37. kill	801dfdcc	T
39. getppid	801dc214	T
41. dup	801cab04	T
42. pipe	801edbe4	T
43. getegid	801dc318	T
46. sigaction	801deee8	T
47. getgid	801dc308	T
48. sigprocmask	801df42c	T
49. getlogin	801dd0e8	T
50. setlogin	801dd160	T
51. acct	801c54ec	T
52. sigpending	801df5d0	T
53. sigaltstack	801dfd10	T
54. ioctl	801ebd1c	T
55. reboot	801e8090	T
56. revoke	800b43f8	T
57. symlink	800b1b58	T

For comparison...
some iOS system
calls

What do we need to make this work?

User program **invokes or calls** subroutine; OS code performs operation; **Returns** control to user program


- The actual code of the service routine
- Mechanism for invocation
 - TRAP Instruction, e.g., TRAP x20
 - TRAP vector
 - $MAR \leftarrow ZEXT[trapvector]$
 - $MDR \leftarrow MEM[MAR]$
 - $PC \leftarrow MDR$
- How to return to user program after execution of OUT completes?

Address	Contents	Comments
x0000		;system space;
x0020	x0400	; Trap vector table
x00FF		; End of trap vector table
x0400		; code for GETC
....		
x0430		; code for OUT
...		
x3000		; user program
...		
	TRAP x20	; call to
xFE00		; Device registers

What do we need to make this work?

User program **invokes or calls** subroutine; OS code performs operation; **Returns** control to user program

- The actual code of the service routine
- Mechanism for invocation
 - TRAP Instruction, e.g., TRAP x20
 - TRAP vector
 - $MAR \leftarrow ZEXT[trapvector]$
 - $MDR \leftarrow MEM[MAR]$
 - $R7 \leftarrow PC$
 - $PC \leftarrow MDR$
- Mechanism for resuming user program
 - $RET \equiv JMP R7$



Address	Contents	Comments
x0000		;system space;
x0020	x0400	; Trap vector table
x00FF		; End of trap vector table
x0400		; code for GETC
....	RET	
x0430		; code for OUT
...		
x3000		; user program
...		
	TRAP x20	; call to
xFE00		; Device registers

Putting it all together: 4 Things make TRAPs work

1. TRAP instruction

- used by program to transfer control to OS subroutines
- 8-bit trap vector names one of the 256 subroutines

2. Trap vector table: stores starting addresses of OS subroutines

- stored at x0000 through x00FF in memory

3. A set of OS subroutines

- part of operating system -- routines start at arbitrary addresses (convention is that system code is below x3000) up to 256 routines

4. A linkage back to the user program (RET)

- want execution of the user program to resume immediately after the TRAP instruction

Subroutines

- Hide details of code and package them with an interface
 - Abstract away details
 - Expose only input/output interface
 - **Arguments** and **return values**
- Why is this a good idea in programming?
 - Reuse; shorter programs
 - Protect; system resources and IP from users
 - Simplify; code comprehension
 - Teamwork; allows multiple developers to work on different pieces; libraries
- TRAPs are examples of OS subroutines

JSR and JSRR and RET

JSR:

R7 <- PC;

PC <- PC + SEXT(PCOffset11);

JSRR:

R7 <- PC;

PC <- BaseR;

RET:

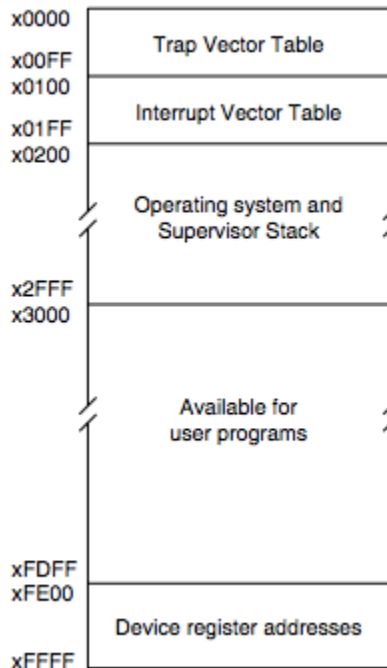
PC <- R7;

Question:

Now that we have TRAPS, and subroutines, can we go back and program using subroutines as planned?

No.

Example of OUT from LC-3



```
.ORIG    x0430

ST R1, SaveR1

Ready    LDI      R1, DSR_ADDR
         BRzp     READY

Write    STI      R0, DDR_ADDR

Return   LD       R1, SaveR1

         RET

DSR_ADDR .FILL    xFE04

DDR_ADDR .FILL    xFE06

SaveR1   .FILL    x0000
```


Observation 2

- When TRAP is called, PC is saved in R7 before service routine starts to execute so that after it finishes, the user program can resume from the right address
- What if R7 had important information? What about other registers that are used internally by the subroutine? (E.g., R1 in OUT)
- Bookmark and more generally bookkeeping!
- We need to save (useful) values in registers in memory before invocation
 - Caller of service routine can save (and restore): **Caller-save**
 - Called service routine saves (and restore): **Callee-save**
- Saving and restoring values of registers is an example of a task computers need to perform in **context switching**

```
; Caller-save user program
...
ST R0, SaveR0      ; store R0 in memory
ST R7, SaveR7      ; store R7 in memory
GETC               ; call TRAP which
                   ; destroys R0 and R7
LD R7, SaveR7      ; restore R7
...               ; consume input in R0
LD R0, SaveR0      ; restore R0
...
HALT

SaveR0 .BLKW 1
SaveR7 .BLKW 1
```

Multiply



- Write an LC-3 assembly language program to calculate:

$$y = ax^2 + bx + c;$$

Summary of Lecture 3

- Need: reuse code (MULT, OUT), hide details from user (xFE02), protect system critical resources (e.g., access to KBDR)
- Idea: package repeated code into **subroutines or system routines**
- Implementation
 - Invocation: TRAP x21 ($R7 \leftarrow PC$, $PC \leftarrow x0430$)
 - System routine code for OUT at x0430 (in OS space) executes
 - RET ($PC \leftarrow R7$); resume execution of user program
- Side effect: May lose useful information in registers (R7, R1)
- Caller save or callee save the relevant registers

JSR and JSRR




TEMP ← PC

If (IR[11] == 0) PC ← BaseR

Else PC ← PC + SEXT(IR[10:0])

R7 ← TEMP



RET ≡ JMP R7

PC ← R7

Examples: a Subtraction subroutine

; SUBTR

A Multiplication subroutine

; IN

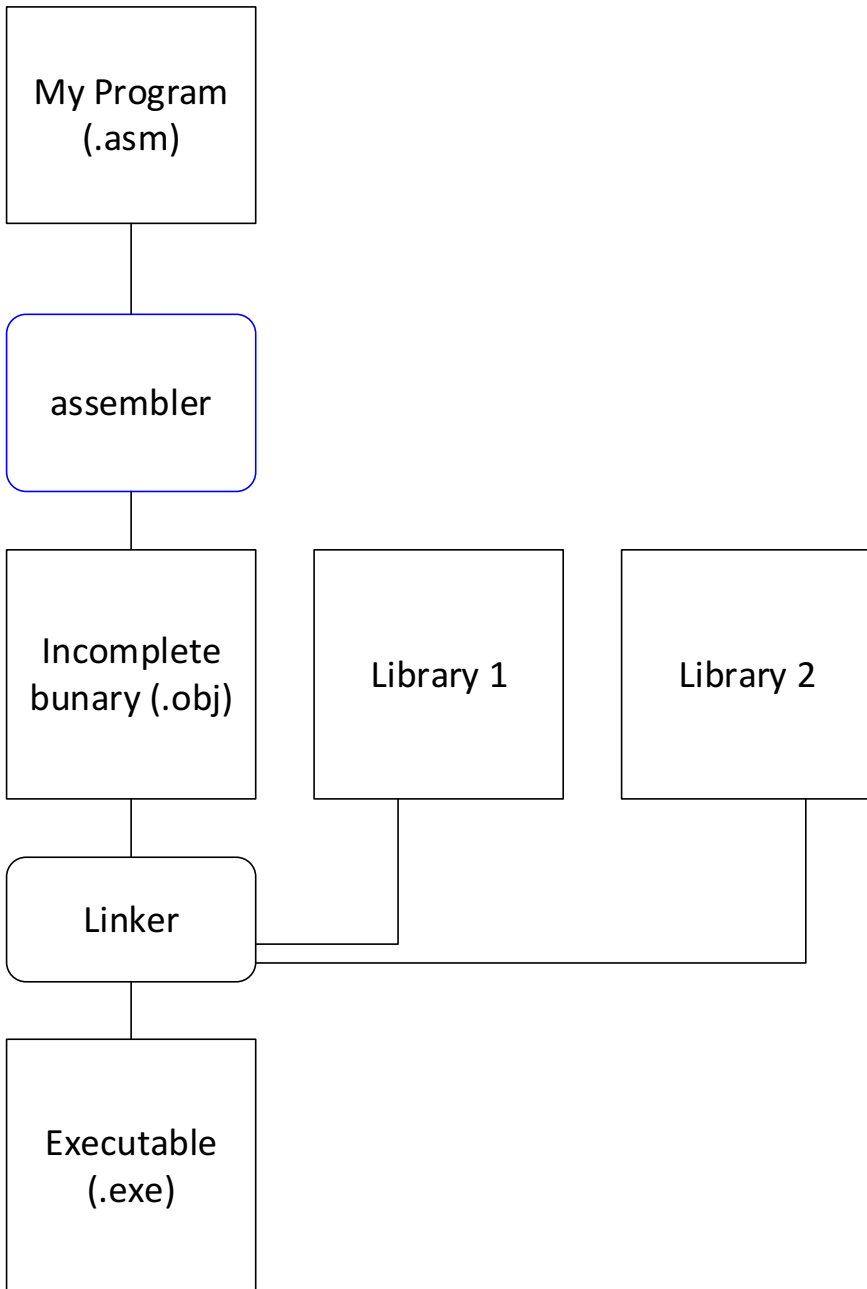
; OUT

;

Subroutine abstraction

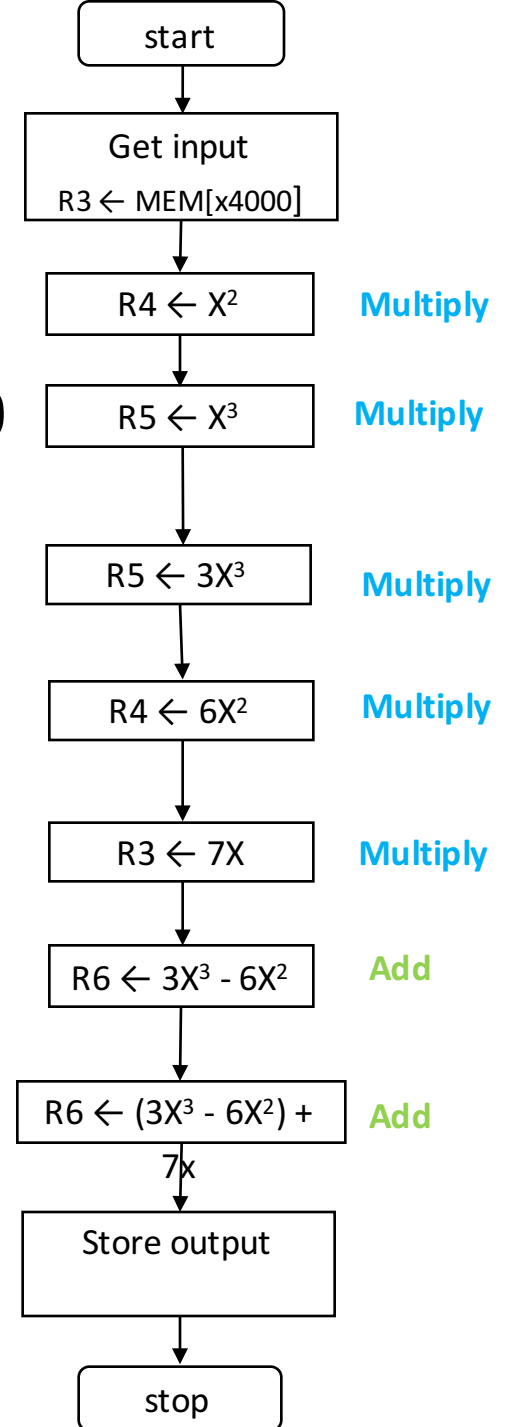
- Interface specifies
 - Input: type (ASCII, int) and location (Registers)
 - Output: type and location
- Optionally used resources (Registers)
- Does not specify?

Libraries and compilation



Exercise

Compute $y=3x^3-6x^2+7x$ for any input $x > 0$



Exercises

- Write an LC-3 assembly language program to calculate: $y = ax^2 + bx + c$;
- Lab exercise: encode a given string by shifting each character by an offset
- Count the number of occurrences of the word “code” in a given string of text.

Nested Subroutines

Can a subroutine call itself?

Summary and tradeoffs

- Idea: package repeated code into subroutines: easier to program, debug, maintain, share
 - TRAP subroutines addressed by their trapvector using the trapvector table
- Jump and return accomplished by setting PC values
- Side effect: May lose useful information in registers
 - Caller-save or callee-save register values
- But, **context switch** (saving and restoring registers) incurs additional cost---cycles, memory accesses
 - Inline functions
- **Next:** stack data structure